

# Why Does ChatGPT “Delve” So Much? Exploring the Sources of Lexical Overrepresentation in Large Language Models

Anonymous COLING 2025 submission

## Abstract

Scientific English is currently undergoing rapid change, with words like “delve,” “intricate,” and “underscore” appearing far more frequently than just a few years ago. It is widely assumed that scientists’ use of large language models (LLMs) is responsible for these trends. We use a novel method to identify 21 focal words whose increased occurrence in scientific abstracts is likely the result of LLM usage. We then pose “the puzzle of lexical overrepresentation”: why are such words overused by LLMs? We fail to find evidence that lexical overrepresentation is caused by model architecture, algorithm choices, or training data. To assess whether reinforcement learning from human feedback (RLHF) contributes to the overuse of focal words, we undertake comparative model testing and conduct an exploratory online study. While the model testing is consistent with RLHF playing a role, our experimental results suggest that participants may be reacting differently to “delve” than to other focal words. We note that lack of transparency remains an obstacle to such research on the sources of lexical overrepresentation.

## 1 Introduction

Like all human language, Scientific English has changed substantially over time (Degaetano-Ortlieb and Teich, 2018; Degaetano-Ortlieb et al., 2018; Bizzoni et al., 2020; Menzel, 2022). New discoveries have fueled (and perhaps been fueled by) the introduction of new lexical items into scientific discourse (Degaetano-Ortlieb and Teich, 2018).<sup>1</sup> Changes in dominant methodological and explanatory frameworks – such as the rise of mechanical philosophy, or the mathematization of scientific fields – have been accompanied by changes in word usage and syntactic structures as well (Degaetano-Ortlieb and Teich, 2018; Krielke, 2024). Such

changes continue through the present (Banks, 2017; Leong, 2020).

Over the last two years, however, Scientific English has witnessed increasing usage of certain lexical items at a seemingly unprecedented pace. Discussions on social media (e.g., Koppenburg 2024; Nguyen 2024) and in academic discourse (Gray, 2024; Kobak et al., 2024; Liang et al., 2024b; Liu and Bu, 2024; Matsui, 2024) have pointed out that words such as “delve,” “intricate,” and “nuanced” have appeared far more frequently in scientific abstracts from 2023 and 2024 compared to earlier years. Unlike many previous changes in Scientific English, these trends do not seem to be explained by changes in the content of science or in wider language use. Instead, it is widely assumed that the sharp increase is due to the use of large language models (LLMs) like ChatGPT for scientific writing. Evidence supporting this hunch has recently emerged (e.g., Cheng et al. 2024; Liang et al. 2024a).

The goals of the present research were twofold. First, we aimed to provide a systematic characterization of this linguistic phenomenon. Some existing work has relied on informal methods to identify words observed to occur more frequently in AI-generated writing (e.g., Matsui 2024). We developed a method for extracting lexical items of interest, described in Section 2, which is rigorous, reproducible, and transferable to other data and models. We identified 21 “focal words”: lexical items that have recently spiked in Scientific English and are overused by ChatGPT-3.5 in scientific writing tasks.

Prior research has focused on quantifying such focal words’ increasing prevalence and estimating how much recent scientific writing has been produced with LLM assistance (e.g., Kobak et al. 2024; Liang et al. 2024b). By contrast, our second goal was to explore the factors that might contribute to the phenomenon of lexical overrepresentation:

<sup>1</sup>Anonymous repo for our code: [github.com/arizus/delve](https://github.com/arizus/delve).

Why does ChatGPT use “delve” (and other focal words) so frequently when generating scientific text? We identified a set of possible factors, characterized in Section 3, and began to assess them. We did not find evidence that model architecture or algorithmic decisions play a major role in the overrepresentation of focal words (Section 5), nor that lexical overrepresentation stems from training or fine-tuning data (Section 4).

LLM training often involves reinforcement learning based on information about quality outputs from human evaluators. We found mixed evidence that reinforcement learning from human feedback (RLHF) contributes to the overrepresentation of our focal words in LLM-generated text. Positive evidence comes from model testing on Meta’s Llama LLM (Section 5). An exploratory experiment described in Section 6 is inconclusive, although our findings indicate that participants became wary of the word “delve” in the first sentence of an abstract (e.g., ‘This article delves into ...’). Since the experiment’s inconclusiveness stems partly from methodological issues, we believe a follow-up study is warranted. Many important questions about the future of LLM-driven language change remain (Section 7).

## 2 Corpus Analysis: Identification of Overrepresented Lexical Items

To probe recent changes in Scientific English, we used PubMed’s publicly available repository of scientific abstracts, which focuses on biomedical literature (National Library of Medicine, 2023) (downloaded through the PubMed API using a Python script (Python Software Foundation, 2024); Snapshot: May 4, 2024; all code on the anon. GitHub). Our analysis includes more than 5.2 billion tokens (inflected forms) from 26.7 million abstracts. To track changes in word usage over time, we measured occurrences per million (opm) of a given token in each year. Figure 1 illustrates the usage trajectories of some baseline items over time. We focus on the period from 1975 to May 2024 as data prior to 1975 are less extensive.

The goal of our corpus analysis was to identify words whose recent overuse in scientific writing is likely the result of LLM deployment. Our approach involved three steps. First, we determined which words were more prevalent in abstracts from 2024 compared to 2020 (since LLMs were not widespread pre-2021). We calculated the percent-




Figure 1: Selected lexical entries: change over time.

age increase in opm for each token in the database between 2020 and 2024. Unsurprisingly, there was a straightforward explanation for why some words spiked in usage during that time. For example, “omicron” and “metaverse” were two of the words that showed the largest percentage increase (see Figure 1). We only considered increases deemed significant by chi-square tests, of which there were about 7300.

We were interested in isolating words whose spike in usage was unexplained. The authors functioned as annotators and independently reviewed the list of words that had the highest percentage change to exclude irrelevant tokens (like year numbers) and words whose spiking had an explanation in terms of scientific advances or world events. In cases of disagreement, we included the word on our list. We stopped once we had 50 words whose usage spiked without any obvious explanation (see incl.ods on GitHub). This list contained several of the words that had been the focus of online conversation, including “delve” and “intricate”.

However, a spike without an obvious explanation is not necessarily LLM-induced. For example, the usage of ‘mash’ increased tenfold, but it is not a word that ChatGPT is known to overuse. The second step of our method involved identifying words that are overrepresented in AI-generated scientific abstracts compared to human-generated abstracts. In producing AI-generated abstracts, our aim was to imitate the process by which researchers might have deployed an LLM in 2022-early 2024 (while paying attention to careful prompt formulation (Wei et al., 2022; Zhou et al., 2022)). After some exploration, we ended up with a two-stage process: (1) We randomly sampled 10,000 abstracts from papers published in 2020 from the PubMed database. Via the API, ChatGPT-3.5 then summarized the associated paper (Prompt: “The following is an abstract of an article. Summarize it in a couple of sentences.”) (2) The ChatGPT-generated summary was then used to ask ChatGPT-3.5 for a

corresponding scientific abstract. (Prompt: “Please write an abstract for a scientific paper, about 200 words in length, based on the following notes.”) We suspect that the most common way of using an LLM to generate an abstract back when ChatGPT could not accept paper-length inputs involved providing important fragments of a paper. We used ChatGPT-3.5 for the entirety of our project because if scientific abstracts in our dataset contain AI-generated language, it is most likely from ChatGPT-3 or ChatGPT-3.5 (Sarkar, 2023).

In total, from 10,000 human abstracts, we generated 9,953 AI abstracts. (For a small number, ChatGPT would not provide a response, presumably due to topic sensitivity.) We then compared the word usage in the AI-generated abstracts with word usage in the original abstracts. We only considered words for which a chi-square test indicated a significant difference in opm between the human- and AI-produced text. This gave us a list of items overused by ChatGPT.

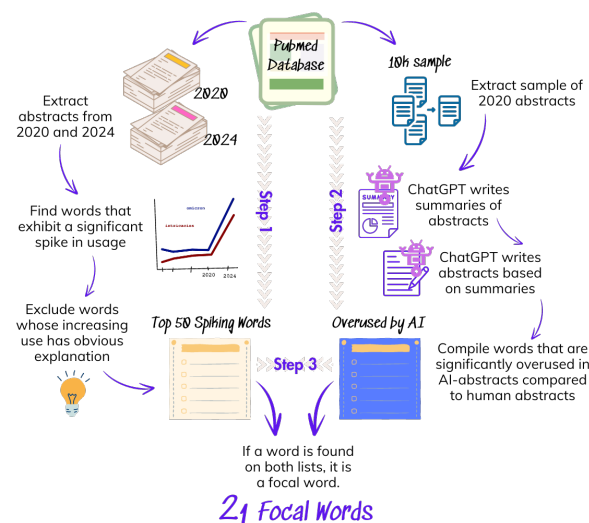


Figure 2: Our method for the systematic identification of focal words.

In the third step of our analysis, we returned to the list of 50 spiking words to ask: Is the word also on the ChatGPT-overuse list? If so, then it became a “focal word” (Figure 2). This gave us a list of 21 focal words (Figure 3 and Appendix A).<sup>2</sup> Each focal word (a) shows a significant spike in opm between 2020 and 2024, (b) its spike lacks an obvious explanation, and (c) ChatGPT tends to use it significantly more than humans when writing

<sup>2</sup>Since the part-of-speech category is not always clear for a given token, the focal word list contains inflected forms instead of lemmata.

scientific abstracts. Thus, a plausible explanation for the increasing prevalence of each focal word in Scientific English is the use of AI.

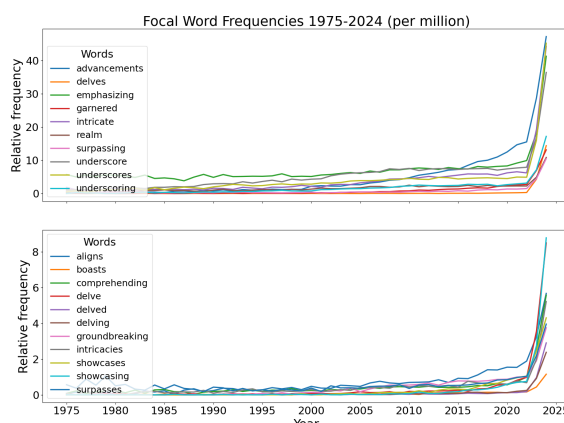


Figure 3: Occurrences per million words in PubMed abstracts for our 21 focal words.

This systematic, three-step method for identifying focal words is novel. It improves on more informal ways of identifying AI-associated words, and it can be applied to other corpora and LLMs beyond ChatGPT-3.5. (Appendix B reports similar results for ChatGPT-4.0(-mini).) Future research can use the method to investigate whether the same words are overrepresented in the outputs of different models – or whether there are LLMs that do not exhibit lexical overrepresentation at all.

### 3 The Puzzle of Lexical Overrepresentation

A question now presents itself: Why are certain words used so often in AI-generated scientific writing? We call this “the puzzle of lexical overrepresentation.” There are a number of factors that might be responsible for the overrepresentation of focal words in scientific abstracts generated by ChatGPT. Importantly, these potential explanations are not mutually exclusive: multiple factors may (and probably do) contribute.

**1. Initial Training Data** Although the focal words are overrepresented relative to human-written abstracts, it is possible that they are not overrepresented relative to the data on which ChatGPT was trained to do next-word prediction. Perhaps these words are actually being used by LLMs with the same frequency as in their training data.

**2. Fine-Tuning Training Data** After LLMs

237	have been trained on next-word prediction,	plain why there is overrepresentation of these	286
238	they are often fine-tuned. For instance, chat-	words in model outputs.	287
239	bots are presented with sample dialogues to		
240	familiarize them with the structure of a con-	7. <b>Other factors</b> This list of potential explana-	288
241	versation. It is possible that something about	tions is not exhaustive. Many other choices	289
242	ChatGPT’s fine-tuning data leads it to favor	– e.g., parameter settings, including tempera-	290
243	certain words (e.g., if the focal words are over-	ture, Top K – might influence lexical overrep-	291
244	represented in the sample dialogues).	resentation in LLM outputs.	292
245			
246	3. <b>Architecture</b> Another possibility is that there	Apportioning responsibility for lexical overrepre-	293
247	is something about the architecture of LLMs,	sentation to these factors is not straightforward.	294
248	or perhaps ChatGPT in particular, that causes	The puzzle of lexical overrepresentation arises in	295
249	them to overuse certain words. Maybe LLMs’	part because LLMs are to a large extent “black	296
250	transformer architecture tends to privilege	boxes” (Knight, 2017; Sculley et al., 2015). Pending	297
251	some lexical items over others in an as-yet-	further advances in LLM explainability or inter-	298
252	unrecognized way. (Even if this explanation	pretability (e.g., Templeton 2024), we will not	299
253	proves correct, the question remains why this	understand many aspects of their behavior. An ad-	300
254	particular set of words is overrepresented.)	ditional obstacle, however, is that many aspects	301
255		of LLM construction are closely-guarded secrets.	302
256	4. <b>Choice of Algorithms</b> LLM development in-	Information that would help discriminate between	303
257	volves many different algorithms. Tokeniza-	the potential explanations above is not public, even	304
258	tion algorithms, for example, segment an in-	for open source models. For instance, we do not	305
259	put string into discrete lexical items. It is	know exactly what data LLMs are trained on (rele-	306
260	possible that the choice of one algorithm over	vant to #1 above), which fine-tuning steps there are	307
261	others causes lexical overrepresentation. Why	(#2), what genres the models are exposed to during	308
262	the algorithm does so, and why these particu-	training (#5), and who the human evaluators are	309
263	lar words are overused, would then be further	(#6). In the remaining sections we pursue several	310
264	questions.	indirect ways of probing potential explanations of	311
265		the puzzle of lexical overrepresentation.	312
266	5. <b>Context Priming</b> A well-known strength of		
267	LLMs is sensitivity to genre. Their outputs	4 <b>Searching for Overrepresentation in</b>	313
268	are highly dependent on the domain and style	<b>Possible Training Data</b>	314
269	requested by the prompt. Perhaps there is		
270	something about being asked to write scien-	Our focal words are overrepresented in text gen-	315
271	tifically that causes ChatGPT to overuse the	erated by ChatGPT compared to earlier PubMed	316
272	focal words. That is, maybe ChatGPT as-	abstracts. Other research indicates that such words	317
273	sociates scientific writing in particular with	also appear less frequently in related datasets in the	318
274	words like “delve” and “intricate.” This expla-	pre-LLM era (Liang et al., 2024b,a; Gray, 2024).	319
275	nation, if correct, raises the further question	Although we do not know exactly what data LLMs	320
276	of why ChatGPT has this association.	have been trained on, these results cast doubt on	321
277		the hypothesis that ChatGPT is using words like	322
278	6. <b>Reinforcement Learning from Human</b>	“delve” and “surpass” frequently because they oc-	323
279	<b>Feedback (RLHF)</b> Human feedback is used	cure frequently in its training data.	324
280	in later training stages to give LLMs informa-	To further demonstrate that the focal words are	325
281	tion about what a quality output looks like. A	probably not overrepresented in the training data,	326
282	human evaluator might rate several potential	we analyzed several additional datasets, namely:	327
283	responses, for example, with the model then	Arxiv abstracts (accessed 4 Aug 2024; contains	328
284	trained with reinforcement learning to pro-	data from 1986 onwards, averaged over all years),	329
285	duce responses similar to highly-rated exem-	the Leipzig Corpus Collective (Goldhahn et al.	330
	plars. It is possible that this human feedback	2012; the English LCC contains mostly news texts	331
	encodes a preference for certain words. If re-	and transcriptions, data from 2005 onwards; pre-	332
	sponses containing “delve” and “intricate” are	processed snapshot from a previous project), and	333
	rated more highly by evaluators, it would ex-	Wikipedia articles and discussions (Foundation	334



2024, accessed 4 Aug 2024). The results are presented in Appendix B. The opm of the focal words in our ChatGPT-3.5-generated abstracts far exceeds their opm in the four datasets examined.

Second, we conducted a similar analysis for various varieties of English using the International Corpus of English (ICE; Kirk and Nelson 2018). Although ICE is relatively small compared to the other datasets (the subcorpora for most varieties contain about one million words), we do not find evidence that the focal words are especially prevalent in any particular variety of English (see Appendix G). This suggests that the overrepresentation of focal words in ChatGPT’s outputs is probably not due to an overrepresentation of a certain variety of English in its training data. It has been hypothesized that LLMs might frequently use words like “delve” because they are more common in varieties of English spoken by human evaluators who provide fine-tuning data, such as Nigerian English (Hern, 2024). Our initial analysis of ICE does not support this hypothesis.

## 5 Model Choices: Architecture and Algorithms

Could choices about model architecture or algorithms be responsible for the puzzle of lexical overrepresentation? To probe this, we would ideally build an LLM ourselves and test the impact of each potential factor on the prevalence of focal words. This requires vast resources, however, and is beyond most researchers’ capabilities, including our own. A more feasible alternative would be to investigate a model that has several released variants – e.g., different versions of the same model using different optimization algorithms. Such a model must also be queryable with respect to information-theoretic measures like entropy (Shannon, 1948). To our knowledge, no LLM offers such fine-grained releases.

The closest we could find is the comparison between Llama 2-Base (Llama-2-7b-hf) and Llama 2-Chat (Llama-2-7b-chat-hf; Touvron et al. 2023). We used the Llama 2 models because they are more similar to ChatGPT-3.5 than Llama 3 (Chiang et al. 2024; but Llama 3 produces similar results; Appendix D). The main difference between these two versions of Llama is that Llama 2-Chat includes fine-tuning and RLHF, whereas Llama 2-Base does not. Llama models can also be queried for per-word entropy (Jurafsky and Martin, 2024).

$$H_{\text{p-word}} = -\frac{1}{L} \sum_{i=1}^n p(x_i) \log p(x_i) \quad (1)$$

By comparing the two models’ per-word entropy for human- and AI-generated abstracts, we could assess which was more “surprised” by abstracts with an overrepresentation of focal words. Any difference between the models provides evidence about the source of lexical overrepresentation. We provided our sample of 10,000 human-written abstracts to both versions of Llama 2, followed by the abstracts rewritten by ChatGPT-3.5 (see Section 2). The results are presented in Table 1.

	Llama 2-Base	Llama 2-Chat
Human	1.616	1.051
AI	1.633	0.886

Table 1: Per-word entropy for human abstracts compared to ChatGPT-generated abstracts. Higher values of entropy mean that the model is more “surprised.”

We observe that Llama 2-Base is slightly less “surprised” by human-written text, while Llama 2-Chat is considerably less “surprised” by AI-generated abstracts, in which the focal words are overrepresented. This suggests the overuse of focal words might be driven by some factor that differs between the models. Given that model architecture and many algorithms are held constant across Llama 2-Base and Llama 2-Chat, our findings suggest that these factors are not the primary causes of lexical overrepresentation. Instead, they indicate that fine-tuning and RLHF – which differ between the models – might be important contributors.

These results are necessarily limited. We cannot claim definitively that the observed difference between the models is driven by the prevalence of focal words rather than some other feature of AI-generated text. Moreover, most of our paper is concerned with ChatGPT rather than Llama. The difficulty is that there are no models of ChatGPT (v.3 or above) that can be queried in the described fashion. We think Llama is a useful approximation.

## 6 RLHF: An Experimental Approach

Our model testing with Llama suggested that RLHF might contribute to lexical overrepresentation. This hypothesis has intuitive plausibility: when human evaluators assess alternative answers to a query, perhaps they are exhibiting a preference for answers containing certain words. Since LLMs are

trained to align their answers with human preferences, they would learn to use those words more frequently (Christiano et al., 2017; Ziegler et al., 2019). To further investigate this potential explanation, we conducted an exploratory online study in which participants indicated whether they preferred scientific abstracts that contained our focal words.

**Materials.** We randomly sampled shorter PubMed abstracts (70-100 words) from the year 2020 and, with Python and using the OpenAI API, used ChatGPT-3.5 to rewrite them with and without focal words. (Shorter abstracts were used to keep stimuli of a manageable length for participants.) For the focal-word abstracts, the prompt included four randomly selected words from our list of 21 focal words. An example prompt is: “Please write a 100-word abstract for the following scientific paper, using words such as ‘delves,’ ‘underscores,’ ‘surpasses,’ and ‘emphasizing’: [SUMMARY].” (The summary was generated via the procedure described in Section 2.) The script instructed ChatGPT to generate and revise an abstract until it contained at least three focal words. For the no-focal-word abstracts, we used a similar prompt: “Please write a 100-word abstract for the following scientific paper, making sure not to use words such as [list of blockwords]: [SUMMARY].” The blockwords included the 21 focal words plus another 21 words identified using the methodology described in Section 2. The script prompted ChatGPT to generate and revise an abstract until it contained none of the blockwords.

We created 200 items, each consisting of one abstract with focal words and one without (for the same paper). We manually filtered out a handful of ungrammatical or nonsensical abstracts. Considerably more than half of the abstracts with focal words included “delve” in the first sentence; we call items containing these abstracts “delve-initial” items. To compile a bank of 30 critical items, we selected the 15 delve-initial items and the 15 other items with the smallest difference in length between the abstracts with and without focal words. (We capped delve-initial items at 50% to prevent participants from detecting the study’s purpose.) We also constructed 30 pairs of distractor items in the same manner as the critical items, except both abstracts were generated using the no-focal-word prompt. A full list of experimental items can be found on Github, and two examples are in Appendix G.

**Participants.** We used Prolific (prolific.com) to

recruit participants. Public information about the human evaluators employed to provide feedback in RLHF is limited (Ouyang et al., 2022; Perrigo, 2023), so we recruited 201 participants from India (140 male, 61 female). Average age was 31.3 years (stdev: 10.6). We also collected data on self-assessed English proficiency and first languages (see anon. GitHub). Participants were compensated at an average rate of \$15 per hour.

**Task and Exclusions.** The study began with IRB information, followed by task instructions, and then the items. An image of the interface can be found in Appendix E. Participants evaluated 20 items in total, indicating which abstract they preferred out of the two presented. The first item was a calibration item, followed by (in random order) 5 critical items, 10 distractor items, 2 items checking language abilities, and 2 attention checks (“This is not a real item, please click on the left button” inserted in the middle of the text). Thus, the proportion of critical items was 25%. Each time an item was displayed, it was randomly determined which abstract was displayed on the left vs. right. If a participant failed one of the attention checks, their data were disregarded. Participants were warned if they were proceeding unrealistically fast ( $0.25 * (225 \text{ ms} + 25\text{ms} * \text{character length of an item})$ ), and items with excessively fast rating times were excluded from our analysis. We also excluded data from participants who completed less than 10 out of the 20 items. After exclusions, we analyzed a total of 1822 ratings, with 1215 ratings for distractor items and 607 ratings for critical items, resulting in each critical item receiving an average of 20.2 ratings (stdev: 3.4). Given the study compensation, the high exclusion rate came as a surprise.

**Analysis.** Our original plan was to test all 30 critical items together in a chi-square analysis against the distractor items (an approximation of random choices), to assess whether participants preferred abstracts containing focal words. These results are reported below. However, during the generation of the abstracts, we noticed the aforementioned excess of delves in the first sentence and split the critical items into delve-initial items and other items. A lower N per condition and a higher-than-expected exclusion rate left us considerably below the originally estimated sample size from a pre-study power analysis. Thus, we added an exploratory mixed-effects logistic regression model, with rating as the dependent variable and condition as the independent variable, including items as a random effect

(rating condition + (1 | item\_id)). Distractor items served as the intercept condition. For delve-initial items and other items, a preference for the focal-word abstract was encoded with 0, and a preference for the no-focal-word abstract with 1. For the distractor items, there are two no-focal-word abstracts, randomly encoded as 0 or 1.

Results. Contrary to our expectations, when all critical items are analyzed together, there is a slight preference for the no-focal-word abstracts. However, this overall difference between all critical items and distractor items is not significant in a chi-square test ( $p = 0.174$ ). The follow-up analysis suggests that this outcome might be driven by the delve-initial items, as Figure 4 illustrates. In the logistic regression model, we observe that the coefficient for the distractor items, represented by the intercept condition, is 0.500 (rounded to the third digit). This indicates that participants did not exhibit a significant preference between the distractor item abstracts, validating our methodology (Appendix F). The analysis also shows that delve-initial items differed significantly from the distractors ( $p = 0.023$ ), with a coefficient of 0.082, indicating that for the delve-initial items, participants preferred the abstracts without focal words. Participants exhibited a slight but non-significant preference for abstracts with focal words for the other critical items (coefficient = -0.017;  $p = 0.651$ ). The group variance was small (0.003), indicating that most of the variability in the ratings was due to the fixed effects. The model converged successfully (log-likelihood = -1324.9522, mean group size = 30.4). A Wald test to determine whether delve-initial items and the other items differed from each other was statistically significant ( $p = 0.03$ , Wald Test Statistic: 4.77).

In looking at the responses for each individual item, we consider a preference for the focal-word or no-focal-word abstract of a given pair to be robust if a random outcome falls outside the margin of error, and marginal otherwise (illustrated for the distractors in Appendix F). This analysis shows a slight difference between delve-initial items and the other critical items: participants exhibit a preference for the no-focal-word abstract in more of the delve-initial items, and a preference for the focal-word abstract in more of the other items.

What explains the difference between delve-initial and the other critical items? We suspect that some participants became or were already sensitive to the occurrence of “delve.” Participants were

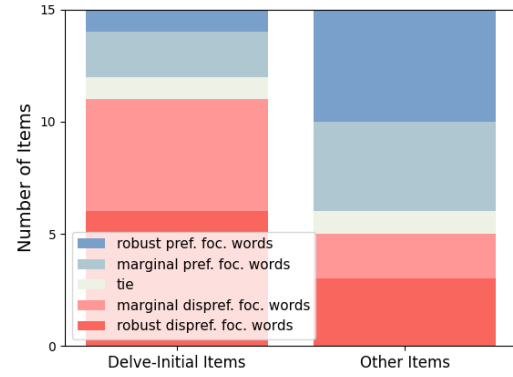


Figure 4: Experimental results: Preferences between focal-word and non-focal-word abstracts in delve-initial and other items.

probably disproportionately young people with an affinity for technology, and so more likely to be familiar with the discourse surrounding AI language use. Wariness about the word “delve” might explain why participants preferred the abstracts without focal words in the delve-initial items, though we would like to see these results confirmed with a larger sample.

Having split the critical items in two, a higher N is needed to draw any conclusions about RLHF as a source of lexical overrepresentation, particularly given that we would expect a preference for focal-word abstracts to be subtle. The study warrants a follow-up. We believe that forcing ChatGPT to use certain words when generating abstracts was suboptimal. For example, if an abstract does not initially convey anything about exceeding or outperforming, then a rewritten abstract that includes the focal word ‘surpasses’ will naturally be worse than the no-focal-word baseline. We suspect that generating critical items in a different way would yield clearer results.

## 7 Discussion and Concluding Remarks

It has been observed that LLMs overuse certain lexical items, a fact even acknowledged by OpenAI (OpenAI, 2024). Our work formalized this finding and identified 21 focal words whose usage has spiked in scientific abstracts and that are overused by ChatGPT-3.5. These results provide additional evidence that recent changes to Scientific English are partly driven by AI. Our work also explored possible explanations of the puzzle of lexical overrepresentation. We failed to find evidence that train-



ing data, model architecture, or algorithm choices play a role. However, model testing with Llama was consistent with the hypothesis that RLHF contributes to overuse of particular words by ChatGPT. Our experimental results suggest that human evaluators may treat “delve” differently from other focal words.

Future research should further probe the impact of each factor canvassed in Section 3 on lexical overrepresentation. (This includes model choices and training data; despite our negative results, we suspect that these factors do influence the lexical choices of LLMs.) We would especially like to see further work on the role of RLHF. Unfortunately, there are several obstacles to such research, particularly the lack of procedural and data transparency surrounding LLM development (Longpre et al., 2024). Moreover, it seems that companies building LLMs often solicit feedback from workers who are underpaid, stressed, and under time pressure (Toxtli et al., 2021; Roberts, 2022; Novick, 2023). It is difficult to simulate these conditions ethically in a research environment. Many online recruitment platforms, including Prolific, rightly require decent compensation.

Although it complicates further study, we think this economic reality lends plausibility to RLHF as a source of lexical overrepresentation. Rushed human evaluators might base their evaluations on the presence of particular words rather than on content, as the former is easier and quicker to evaluate than the latter. If certain words are treated as a proxy for quality, that could explain their overrepresentation in LLM outputs. This mechanism coheres with our impression that a major social consequence of LLMs is the decoupling of form and content. Many of us take fluency or style as a signal of quality content (McNamara et al. 2010, and in an L2 context Kim and Crossley 2018). Because LLMs are masterful at generating fluid text in just about any style, this heuristic is radically undermined by the increasing ubiquity of LLM-generated text. The irony is that, if our hypothesis about RLHF proves correct, this heuristic has shaped model training as well. LLMs may be undercutting the very same heuristic that has shaped their own lexical preferences.

In addition to further exploring potential explanations of the puzzle of lexical overrepresentation, it would be interesting to apply the present methods to LLMs besides ChatGPT. Our impression is that ChatGPT and Llama overuse many of the same

words, but a systematic investigation is needed. This research could also be extended beyond the word level (Ortmann et al., 2021) and to domains beyond Scientific English. We suspect that Scientific English played a minor role in the training of LLMs. It seems more likely that human evaluators rated academic writing in general, with their preferences shaping LLMs’ scientific writing through overspill.

In general, we believe more attention should be paid to how LLMs are changing language. Almost all of our 21 focal words were already increasing in usage in the years leading up to the release of ChatGPT, suggesting that LLMs may accelerate language change (Matsui, 2024). With the increasing prevalence of AI-generated text in many areas of life, LLMs are arguably influencing the language usage even of people who do not themselves interact with these models. Our findings also show that lexical overrepresentation remains a feature of current iterations of ChatGPT (Appendix B), indicating that the phenomenon is here to stay.

Still, it is difficult to predict just how AI will shape language in the future. Discussions on social media and in academic discourse, plus our exploratory findings for items with “delve,” indicate that there is some public awareness of LLMs’ overuse of particular words. This awareness could influence future rounds of RLHF, leading to a realignment of AI and human preferences. At the same time, the language of today – lexical overrepresentations and all – will become the training data for the models of tomorrow, raising concerns about model degradation over time (Alemohammad et al., 2023; Briesch et al., 2023; Hataya et al., 2023; Shumailov et al., 2023).

One thing is certain: through LLMs, tech companies are having a global impact on language usage. We believe this strengthens the case for broader societal debate about the power and responsibilities of these companies. Moreover, our speculations about how the feedback of rushed and underpaid workers might contribute to lexical overrepresentation compound ethical worries about the poor working conditions of tech companies’ employees in the Global South (Kwet, 2019; Gray, 2024; Rohde et al., 2024). There are thus both moral and non-moral reasons to apply greater scrutiny to how human feedback is collected and used in the training of LLMs.



## References

- Sina Alemohammad, Josue Casco-Rodriguez, Lorenzo Luzi, Ahmed Imtiaz Humayun, Hossein Babaei, Daniel LeJeune, Ali Siahkoohi, and Richard G Baraniuk. 2023. Self-consuming generative models go mad. *arXiv preprint arXiv:2307.01850*.
- David Banks. 2017. The extent to which the passive voice is used in the scientific journal article, 1985–2015, *functional linguistics*, 4 (12), 2–17.
- Yuri Bizzoni, Stefania Degaetano-Ortlieb, Peter Fankhauser, and Elke Teich. 2020. Linguistic variation and change in 250 years of english scientific writing: A data-driven approach. *Frontiers in Artificial Intelligence*, 3:73.
- Martin Briesch, Dominik Sobania, and Franz Rothlauf. 2023. Large language models suffer from their own output: An analysis of the self-consuming training loop. *arXiv preprint arXiv:2311.16822*.
- Huzi Cheng, Bin Sheng, Aaron Lee, Varun Chaudhary, Atanas G Atanasov, Nan Liu, Yue Qiu, Tien Yin Wong, Yih-Chung Tham, and Ying-Feng Zheng. 2024. Have ai-generated texts from llm infiltrated the realm of scientific writing? a large-scale analysis of preprint platforms. *bioRxiv*, pages 2024–03.
- Wei-Lin Chiang, Lianmin Zheng, Ying Sheng, Anastasios Nikolas Angelopoulos, Tianle Li, Dacheng Li, Hao Zhang, Banghua Zhu, Michael Jordan, Joseph E. Gonzalez, and Ion Stoica. 2024. *Chatbot arena: An open platform for evaluating llms by human preference*. Preprint, arXiv:2403.04132.
- Paul F Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. 2017. Deep reinforcement learning from human preferences. *Advances in neural information processing systems*, 30.
- Stefania Degaetano-Ortlieb, Hannah Kermes, Ashraf Khamis, and Elke Teich. 2018. An information-theoretic approach to modeling diachronic change in scientific english. In *From data to evidence in English language research*, pages 258–281. Brill.
- Stefania Degaetano-Ortlieb and Elke Teich. 2018. Using relative entropy for detection and analysis of periods of diachronic linguistic change. In *Proceedings of the second joint SIGHUM workshop on computational linguistics for cultural heritage, social sciences, humanities and literature*, pages 22–33.
- Wikimedia Foundation. 2024. *Wikipedia dump*. Accessed: 4 August 2024.
- Dirk Goldhahn, Thomas Eckart, Uwe Quasthoff, et al. 2012. Building large monolingual dictionaries at the leipzig corpora collection: From 100 to 200 languages. In *LREC*, volume 29, pages 31–43.
- Andrew Gray. 2024. Chatgpt" contamination": estimating the prevalence of llms in the scholarly literature. *arXiv preprint arXiv:2403.16887*.
- Ryuichiro Hataya, Han Bao, and Hiromi Arai. 2023. Will large-scale generative models corrupt future datasets? In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 20555–20565.
- Alex Hern. 2024. *TechScape: How cheap, outsourced labour in Africa is shaping AI English*. Accessed: 2024-08-12.
- Dan Jurafsky and James H. Martin. 2024. *Speech and Language Processing*. Online draft. 3rd ed. draft, Feb 3, 2024 release.
- Minkyung Kim and Scott A Crossley. 2018. Modeling second language writing quality: A structural equation investigation of lexical, syntactic, and cohesive features in source-based and independent writing. *Assessing Writing*, 37:39–56.
- John Kirk and Gerald Nelson. 2018. The international corpus of english project: A progress report. *World Englishes*, 37(4):697–716.
- Will Knight. 2017. The dark secret at the heart of ai.
- Dmitry Kobak, Rita González Márquez, Emőke-Ágnes Horvát, and Jan Lause. 2024. Delving into chatgpt usage in academic writing through excess vocabulary. *arXiv preprint arXiv:2406.07016*.
- Patrick Koppenburg. 2024. Tweet on 01 april 2024. <https://x.com/PKoppenburg/status/1774757167045788010>. Accessed: 2024-08-12.
- Marie-Pauline Krielke. 2024. Cross-linguistic dependency length minimization in scientific language: Syntactic complexity reduction in english and german in the late modern period. *Languages in Contrast*, 24(1):133–163.
- Michael Kwet. 2019. Digital colonialism: Us empire and the new imperialism in the global south. *Race & Class*, 60(4):3–26.
- Alvin Ping Leong. 2020. The passive voice in scientific writing through the ages: A diachronic study. *Text & Talk*, 40(4):467–489.
- Weixin Liang, Zachary Izzo, Yaohui Zhang, Haley Lepp, Hancheng Cao, Xuandong Zhao, Lingjiao Chen, Hao-tian Ye, Sheng Liu, Zhi Huang, et al. 2024a. Monitoring ai-modified content at scale: A case study on the impact of chatgpt on ai conference peer reviews. *arXiv preprint arXiv:2403.07183*.
- Weixin Liang, Yaohui Zhang, Zhengxuan Wu, Haley Lepp, Wenlong Ji, Xuandong Zhao, Hancheng Cao, Sheng Liu, Siyu He, Zhi Huang, et al. 2024b. Mapping the increasing use of llms in scientific papers. *arXiv preprint arXiv:2404.01268*.
- Jialin Liu and Yi Bu. 2024. Towards the relationship between aigc in manuscript writing and author profiles: evidence from preprints in llms. *arXiv preprint arXiv:2404.15799*.

823	Shayne Longpre, Robert Mahari, Anthony Chen, Naana Obeng-Marnu, Damien Sileo, William Brannon, Niklas Muennighoff, Nathan Khazam, Jad Kabbara, Kartik Perisetla, et al. 2024. A large-scale audit of dataset licensing and attribution in ai. <i>Nature Machine Intelligence</i> , 6(8):975–987.	874
824		875
825		876
826		877
827		
828		
829	Kentaro Matsui. 2024. Delving into pubmed records: Some terms in medical writing have drastically changed after the arrival of chatgpt. <i>medRxiv</i> , pages 2024–05.	878
830		879
831		880
832		881
833	Danielle S McNamara, Scott A Crossley, and Philip M McCarthy. 2010. Linguistic features of writing quality. <i>Written communication</i> , 27(1):57–86.	882
834		883
835		884
836	Katrin Menzel. 2022. Medical discourse in late modern english: Insights from a multidisciplinary corpus of scientific journal articles. In <i>Corpus pragmatic studies on the history of medical discourse</i> , pages 79–104. John Benjamins.	885
837		886
838		887
839		888
840		889
841	National Library of Medicine. 2023. PubMed Database. <a href="https://pubmed.ncbi.nlm.nih.gov/">https://pubmed.ncbi.nlm.nih.gov/</a> . Accessed: 2024-08-12.	890
842		891
843		892
844	Jeremy Nguyen. 2024. Tweet on 30 march 2024. <a href="https://x.com/JeremyNguyenPhD/status/1774021645709295840">https://x.com/JeremyNguyenPhD/status/1774021645709295840</a> . Accessed: 2024-08-12.	893
845		894
846		895
847	Michael Novick. 2023. <i>A.i.’s dirty secret: It’s powered by digital sweatshops</i> . Blog post.	896
848		897
849	OpenAI. 2024. Tweet on 08 april 2024. <a href="https://x.com/ChatGPTapp/status/1777221658807521695">https://x.com/ChatGPTapp/status/1777221658807521695</a> . Accessed: 2024-08-12.	898
850		899
851		900
852	Katrin Ortmann, Adam Roussel, and Stefanie Dipper. 2021. <i>Computational Historical Linguistics: Annotations, Tools &amp; Corpora</i> . Ruhr-Universität Bochum.	901
853		902
854		903
855	Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. 2022. Training language models to follow instructions with human feedback. <i>Advances in neural information processing systems</i> , 35:27730–27744.	904
856		905
857		906
858		907
859		908
860		909
861	Billy Perrigo. 2023. Exclusive: Openai used kenyan workers on less than \$2 per hour to make chatgpt less toxic. <i>Time Magazine</i> , 18:2023.	910
862		911
863		912
864	Python Software Foundation. 2024. <i>Python 3</i> .	913
865	Jennafer Roberts. 2022. <i>The precarious human work behind ai</i> . Blog post.	914
866		915
867	Friederike Rohde, Josephin Wagner, Andreas Meyer, Philipp Reinhard, Marcus Voss, Ulrich Petschow, and Anne Mollen. 2024. Broadening the perspective for sustainable artificial intelligence: sustainability criteria and indicators for artificial intelligence systems. <i>Current Opinion in Environmental Sustainability</i> , 66:101411.	916
868		917
869		918
870		919
871		920
872		
873		
	Sujan Sarkar. 2023. AI Industry Analysis: 50 Most Visited AI Tools and Their 24B+ Traffic Behavior. <a href="https://writerbuddy.ai/blog/ai-industry-analysis">https://writerbuddy.ai/blog/ai-industry-analysis</a> . Accessed: 2024-08-12.	
	David Sculley, Gary Holt, Daniel Golovin, Eugene Davydov, Todd Phillips, Dietmar Ebner, Vinay Chaudhary, Michael Young, Jean-Francois Crespo, and Dan Dennison. 2015. Hidden technical debt in machine learning systems. <i>Advances in neural information processing systems</i> , 28.	
	Claude Elwood Shannon. 1948. A mathematical theory of communication. <i>The Bell system technical journal</i> , 27(3):379–423.	
	Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Yarin Gal, Nicolas Papernot, and Ross Anderson. 2023. The curse of recursion: Training on generated data makes models forget. <i>arXiv preprint arXiv:2305.17493</i> .	
	Adly Templeton. 2024. <i>Scaling monosemanticity: Extracting interpretable features from claude 3 sonnet</i> . Anthropic.	
	Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. 2023. Llama 2: Open foundation and fine-tuned chat models. <i>arXiv preprint arXiv:2307.09288</i> .	
	Carlos Toxtli, Siddharth Suri, and Saiph Savage. 2021. Quantifying the invisible labor in crowd work. <i>Proceedings of the ACM on human-computer interaction</i> , 5(CSCW2):1–26.	
	Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. 2022. Chain-of-thought prompting elicits reasoning in large language models. <i>Advances in neural information processing systems</i> , 35:24824–24837.	
	Denny Zhou, Nathanael Schärli, Le Hou, Jason Wei, Nathan Scales, Xuezhi Wang, Dale Schuurmans, Claire Cui, Olivier Bousquet, Quoc Le, et al. 2022. Least-to-most prompting enables complex reasoning in large language models. <i>arXiv preprint arXiv:2205.10625</i> .	
	Daniel M Ziegler, Nisan Stiennon, Jeffrey Wu, Tom B Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. 2019. Fine-tuning language models from human preferences. <i>arXiv preprint arXiv:1909.08593</i> .	

## A List Of Focal Words

Word	opm 2020	opm 2024	Incr. %
delves	0.21	14.38	6697.14
delved	0.12	2.90	2240.47
delving	0.12	2.38	1816.83
showcasing	0.59	8.79	1396.03
delve	0.58	8.50	1374.92
boasts	0.11	1.15	918.18
underscores	4.50	45.19	903.61
comprehending	0.56	5.58	898.95
intricacies	0.60	5.22	772.85
surpassing	1.37	10.50	667.48
intricate	6.22	44.22	611.24
underscoring	2.70	17.17	536.94
garnered	2.44	13.13	437.19
showcases	0.82	4.31	422.45
emphasizing	8.30	41.27	397.12
underscore	7.42	36.40	390.65
realm	2.25	10.85	381.10
surpasses	0.85	3.96	367.55
groundbreaking	0.87	3.75	330.42
advancements	12.49	47.17	277.59
aligns	1.55	5.68	266.97

Table 2: Our 21 focal words.

## B Analysis Of Further Corpora and GPT-4o

We used the same summaries from the sample of 10,000 abstracts and used a Python script to generate abstracts using GPT-4o-mini, which were then analyzed as per Section 2. The system’s role was set as ‘You are a world-leading scientist.’ and the prompt was the same as that used with GPT-3.5. The results are very similar, with a few exceptions: ‘boasts’ is no longer overused; ‘delve’ is still overused, but to a lesser extent; and the usage of ‘underscore’ has increased significantly. These differences could be artifacts of the methodology (the GPT-4-generated abstracts are based on the same GPT-3.5 summaries used in Section 2), the consequence of active intervention, RLHF workers responding to overuse, and/or other factors. We also conducted a spot-check with a prompt specifying the role ‘You are a helpful assistant.’ using GPT-4o. For each role, we generated 500 abstracts and analyzed them. There was no noticeable difference with GPT-4o-mini.

## C Examples Of Critical Items

### C.1 A delve-initial item

**Focal-word abstract:** This study delves into the impacts of maintaining mean arterial blood pressure (MABP) at a specific level during the prehospital and initial hospital phases of treatment for patients with traumatic spinal cord injury (TSCI). The results show a strong correlation between maintaining MABP at 80 mm Hg and improved long-term neurological outcomes. Specifically, prehospital and operating room levels had the greatest effect, while maintaining MABP at 80 mm Hg during the first 2 days in the neurointensive care unit (NICU) was also beneficial. These findings surpass previous knowledge and highlight the advancements in comprehending the relationship between MABP and TSCI recovery.

**No-focal-word abstract:** This study explores the impact of maintaining mean arterial blood pressure (MABP) at 80mm Hg during prehospital and initial hospital treatment on long-term neurological outcomes in patients with TSCI. Results showed a significant correlation between higher MABP levels and improved outcomes, with the strongest impact observed in the prehospital and operating room phases. The benefits of maintaining MABP at 80mm Hg were also observed in the first 2 days in the NICU. These findings highlight the crucial role of MABP management in minimizing neurogenic shock-induced damage and emphasize the importance of maintaining adequate blood pressure in TSCI patients.

### C.2 A non delve-initial item

**Focal-word abstract:** This paper showcases a novel approach for targeting and disrupting c-di-GMP signaling pathways in bacteria. By utilizing a c-di-GMP-sequestering peptide (CSP), the researchers have developed a method to bind and inhibit c-di-GMP, a key bacterial second messenger. Through structure-based mutations, a more powerful and compact variant of the CSP has been created, effectively preventing biofilm formation in *Pseudomonas aeruginosa*. This advancement holds promise for controlling bacterial behaviors mediated by c-di-GMP and could have implications for the development of new antibacterial strategies. The results of this study highlight the potential of CSP as a tool for delving into the intricate mechanisms of c-di-GMP signaling.

**No-focal-word abstract:** A novel approach

Word	ChatGPT 3.5	ChatGPT 4o-mini	Arxiv	LCC	Pubmed	Wiki
of	45624.84	42622.65	42842.72	27363.47	38634.99	23116.18
and	38889.24	32537.79	26395.28	28488.53	39469.96	21149.63
the	63174.05	55111.23	72009.63	59324.62	52139.05	53379.32
data	978.91	1075.59	2484.20	418.29	1734.75	142.81
results	4074.64	3307.32	2352.13	244.52	1722.07	95.37
i	32.21	61.17	414.03	4715.42	214.82	8041.61
year	78.50	61.77	37.58	1076.29	217.25	397.61
patients	4416.82	3936.56	48.97	131.48	4775.73	23.04
advancements	319.37	407.59	22.54	2.56	15.53	1.11
aligns	6.71	19.99	6.68	1.32	1.89	0.90
boasts	5.37	0.61	0.43	14.11	0.16	1.48
comprehending	6.71	7.27	1.77	0.37	0.99	0.31
delve	19.46	18.17	4.07	2.23	0.98	1.21
delves	183.17	23.01	3.20	0.79	0.32	0.53
delved	6.71	0.61	0.30	0.61	0.18	0.38
delving	8.72	0.61	0.72	0.76	0.24	0.61
emphasizing	138.21	367.61	10.21	2.82	9.92	2.64
garnered	20.80	173.21	4.09	4.34	2.74	4.61
groundbreaking	38.92	17.56	2.47	5.91	1.02	2.26
intricate	163.04	316.14	17.87	4.79	6.22	2.13
intricacies	15.43	27.25	1.98	1.24	0.68	0.68
realm	10.74	54.51	11.53	9.22	2.27	8.46
showcases	28.85	4.24	3.19	4.65	1.05	1.46
showcasing	30.19	58.14	5.89	5.42	0.75	1.65
surpasses	4.03	4.24	11.16	1.14	1.04	0.40
surpassing	5.37	17.56	7.61	1.66	1.51	1.42
underscore	18.12	1365.08	5.17	1.53	7.91	0.72
underscores	60.39	1048.94	4.95	1.90	4.91	0.90
underscoring	10.06	313.71	2.57	0.66	3.15	0.20

Table 3: Occurrences per million for selected baseline words and our 21 focal words. Results are averaged across all given years of the corpus.

has been devised for blocking c-di-GMP signaling pathways, a crucial mechanism in bacterial cell functioning. The technique employs a c-di-GMP-sequestering peptide (CSP) that exhibits strong affinity for c-di-GMP and effectively inhibits its signaling. Through targeted mutations, a potent, shortened variant of CSP has been developed, demonstrating efficient inhibition of biofilm formation in *Pseudomonas aeruginosa*. This innovative method provides a highly promising strategy for targeting c-di-GMP and holds potential for combating various bacterial infections. Further studies could focus on developing more potent and specific CSP variants to fully comprehend and utilize the role of c-di-GMP in regulating bacterial functions.

## D Per-word Entropy for various Llama Models

We validate our results for various Llama models. We use the latest versions available on 28 August 2024. We intend to extend our analysis to the larger, 70bn parameter models. However, due to quota restrictions, we are unable to perform these calculations at this time. We expect that the results will be similar and plan to include them once the account limitations are resolved.

All models show a drop in average per-word entropy for human input when comparing base models to chat models, with a more pronounced drop observed for AI input. Most models show lower entropy values for human text with the base model compared to AI text. This pattern reverses in the



1025 chat models, where AI text shows lower entropy  
1026 values than human text.

8b	Llama 3-Base	Llama 3-Chat
Human	1.862	1.174
AI	1.928	1.165

8b	Llama 3.1-Base	Llama 3.1-Chat
Human	1.854	1.731
AI	1.838	1.653

1027

## E The Rating Interface

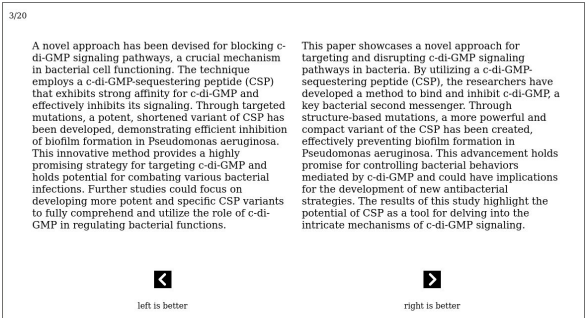


Figure 5: The rating interface for our experiment.

1028

## F Ratings For The Distractor Items




Figure 6: The experimental results for individual distractor items.

1029

1030

## G Analysis of the International Corpus of English

1031 An analysis of the Englishes of the world can be  
1032 found in Figure 7.




Figure 7: Word frequencies for selected lexical items across various English variants.